

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
import sklearn
from sklearn.preprocessing import StandardScaler
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import classification_report
from sklearn.model_selection import RepeatedKFold, cross_val_score

import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
```

```
df=pd.read_csv('dataset_med.csv')
```

```
df.head()
```



	id	age	gender	country	diagnosis_date	cancer_stage	family_history
0	1	64.0	Male	Sweden	2016-04-05	Stage I	Yes
1	2	50.0	Female	Netherlands	2023-04-20	Stage III	Yes
2	3	65.0	Female	Hungary	2023-04-05	Stage III	Yes
3	4	51.0	Female	Belgium	2016-02-05	Stage I	No
4	5	37.0	Male	Luxembourg	2023-11-29	Stage I	No

```
df.shape
```



```
(80733, 17)
```

```
df.columns
```

```
Index(['id', 'age', 'gender', 'country', 'diagnosis_date', 'cancer_stage',
      'family_history', 'smoking_status', 'bmi', 'cholesterol_level',
      'hypertension', 'asthma', 'cirrhosis', 'other_cancer',
      'treatment_type',
      'end_treatment_date', 'survived'],
      dtype='object')
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 80733 entries, 0 to 80732
Data columns (total 17 columns):
#   Column                      Non-Null Count  Dtype
---  -
0   id                          80733 non-null  int64
1   age                         80733 non-null  float64
2   gender                      80733 non-null  object
3   country                     80733 non-null  object
4   diagnosis_date              80733 non-null  object
5   cancer_stage                80733 non-null  object
6   family_history              80732 non-null  object
7   smoking_status              80732 non-null  object
8   bmi                         80732 non-null  float64
9   cholesterol_level           80732 non-null  float64
10  hypertension                 80732 non-null  float64
11  asthma                       80732 non-null  float64
12  cirrhosis                   80732 non-null  float64
13  other_cancer                 80732 non-null  float64
14  treatment_type              80732 non-null  object
15  end_treatment_date          80732 non-null  object
16  survived                    80732 non-null  float64
dtypes: float64(8), int64(1), object(8)
memory usage: 10.5+ MB
```

```
df = df.drop('id', axis=1)
```

```
print(df.duplicated().sum())
```

```
0
```

```
df.describe()
```



	age	bmi	cholesterol_level	hypertension	asthma
count	80733.000000	80732.000000	80732.000000	80732.000000	80732.000000
mean	54.961651	30.496871	233.716791	0.752415	0.466655
std	9.981384	8.376907	43.495299	0.431612	0.498890
min	15.000000	16.000000	150.000000	0.000000	0.000000
25%	48.000000	23.200000	197.000000	1.000000	0.000000
50%	55.000000	30.500000	242.000000	1.000000	0.000000
75%	62.000000	37.800000	271.000000	1.000000	1.000000
max	101.000000	45.000000	300.000000	1.000000	1.000000

```
binary_col = ["hypertension", "asthma", "cirrhosis", "other_cancer", "survived"]
category_col = ["cancer_stage", "smoking_status", "treatment_type"]
```

```
# Fill missing values in binary columns with 0
df[binary_col] = df[binary_col].fillna(0)
```

```
# Fill missing values in numerical columns with the mean
numerical_cols_to_fill = ['age', 'bmi', 'cholesterol_level']
df[numerical_cols_to_fill] = df[numerical_cols_to_fill].fillna(df[numerical_col
```

```
# Fill missing values in categorical columns with a placeholder
df[category_col] = df[category_col].fillna('Unknown')
```

```
df['age'] = df['age'].astype("int8")
df[binary_col] = df[binary_col].astype("int8")
df["gender"] = df["gender"].map({'Female': 1, 'Male': 0}).fillna(0).astype("int
df["family_history"] = df["family_history"].map({'Yes': 1, 'No': 0}).fillna(0).
df[category_col] = df[category_col].astype("category")
```

```
# Convert date columns to datetime objects and calculate treatment_time
df['diagnosis_date'] = pd.to_datetime(df['diagnosis_date'], errors='coerce')
df['end_treatment_date'] = pd.to_datetime(df['end_treatment_date'], errors='coe
df['treatment_time'] = (df['end_treatment_date'] - df['diagnosis_date']).dt.day
```

```
# Fill NaN/NaT values in treatment_time with a placeholder (e.g., -1) before cc
df['treatment_time'] = df['treatment_time'].fillna(-1).astype(int)
```

```
# Drop the original date columns
```

```
df = df.drop(['diagnosis_date','end_treatment_date'],axis=1)
```

```
df.info()
```

```
df.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 80733 entries, 0 to 80732
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    80733 non-null  int8
1   gender                 80733 non-null  int8
2   country                80733 non-null  object
3   cancer_stage           80733 non-null  category
4   family_history         80733 non-null  int8
5   smoking_status        80733 non-null  category
6   bmi                    80733 non-null  float64
7   cholesterol_level     80733 non-null  float64
8   hypertension           80733 non-null  int8
9   asthma                 80733 non-null  int8
10  cirrhosis              80733 non-null  int8
11  other_cancer           80733 non-null  int8
12  treatment_type         80733 non-null  category
13  survived               80733 non-null  int8
14  treatment_time         80733 non-null  int64
dtypes: category(3), float64(2), int64(1), int8(8), object(1)
memory usage: 3.3+ MB
```

	age	gender	country	cancer_stage	family_history	smoking_status	bmi
0	64	0	Sweden	Stage I	1	Passive Smoker	29.4
1	50	1	Netherlands	Stage III	1	Passive Smoker	41.2
2	65	1	Hungary	Stage III	1	Former Smoker	44.0
3	51	1	Belgium	Stage I	0	Passive Smoker	43.0
4	37	0	Luxembourg	Stage I	0	Passive Smoker	19.7

```
from sklearn.preprocessing import LabelEncoder
import pandas as pd

le = LabelEncoder()
df['cancer_stage'] = le.fit_transform(df['cancer_stage'])

df = pd.get_dummies(df, columns=['smoking_status', 'treatment_type'], drop_first=True)

df.head()
```



	age	gender	country	cancer_stage	family_history	bmi	cholesterol_level
0	64	0	Sweden	1	1	29.4	190
1	50	1	Netherlands	3	1	41.2	280
2	65	1	Hungary	3	1	44.0	260
3	51	1	Belgium	1	0	43.0	240
4	37	0	Luxembourg	1	0	19.7	170

5 rows × 21 columns


```
fig, axes = plt.subplots(1, 3, figsize=(18, 5))

sns.histplot(df, x="age", binwidth=3, edgecolor="white", element="bars", kde=True,
             axes[0].set_title("Age of patients"))

sns.histplot(df["bmi"], bins=30, edgecolor="white", kde=True, ax=axes[1])
axes[1].set_title("BMI Distribution of patients")

sns.histplot(df["cholesterol_level"], bins=30, edgecolor="white", kde=True, ax=
axes[2].set_title("Cholestrol Level of patients"))

plt.tight_layout()
fig.show()
```




```
cols = ["hypertension", "asthma", "cirrhosis", "other_cancer", "survived"]

fig, axes = plt.subplots(2, 3, figsize=(18, 10))

axes = axes.flatten()

for i, col in enumerate(cols):
    sns.countplot(x=col, data=df, ax=axes[i], color="#4682A9")
    for p in axes[i].patches:
        height = p.get_height()
        axes[i].annotate(f'{int(height)}',
                        (p.get_x() + p.get_width() / 2., height),
                        ha='center', va='bottom')
    axes[i].set_title(f"Distribution of {col.capitalize()}")
    axes[i].set_xlabel(col.capitalize())
    axes[i].set_ylabel("Count")
```

```
if len(cols) < len(axes):  
    fig.delaxes(axes[-1]) # Remove unused subplot  
  
plt.tight_layout()  
plt.show()
```



```

category = ['gender', 'cancer_stage',
            'family_history']

fig, axes = plt.subplots(1, 3, figsize=(18, 5))

for ax, col in zip(axes.flatten(), category):
    sns.countplot(data=df, x=col, ax=ax, color="#4682A9")


    for p in ax.patches:
        height = p.get_height()
        ax.annotate(f'{int(height)}',
                    (p.get_x() + p.get_width() / 2, height),
                    ha='center', va='bottom', fontsize=9)

    ax.set_title(f'Count of {col}')
    ax.tick_params(axis='x', rotation=45)

axes_flat = axes.flatten()
if len(category) < len(axes_flat):
    for i in range(len(category), len(axes_flat)):
        fig.delaxes(axes_flat[i])

plt.tight_layout()
plt.show()

```






```
plt.figure(figsize=(18, 5))
ax = sns.countplot(data=df, x="country")
plt.title("Count of country")
plt.xticks(rotation=45)

# Add labels on top of bars
for p in ax.patches:
    height = p.get_height()
    ax.annotate(f'{int(height)}',
                (p.get_x() + p.get_width() / 2, height),
                ha='center', va='bottom')

plt.tight_layout()
plt.show()
```



```
df_model=df.copy()
```




5 rows x 21 columns

Page 10 of 17

family_history	1.000000	-1.065642e-03
bmi	-0.001066	1.000000e+00
cholesterol_level	-0.000358	7.480522e-01
hypertension	-0.004445	2.758397e-03
asthma	-0.005626	-3.059462e-04
cirrhosis	0.003925	1.329084e-03
other_cancer	-0.000724	-2.268058e-03
survived	-0.004931	-4.734229e-03
treatment_time	-0.003118	-4.295175e-03
smoking_status_Former Smoker	0.003670	-4.723461e-03
smoking_status_Never Smoked	0.001097	6.496740e-04
smoking_status_Passive Smoker	-0.005782	4.918737e-04
smoking_status_Unknown	-0.003513	-9.354595e-19
treatment_type_Combined	0.001464	3.749931e-03
treatment_type_Radiation	-0.004656	-6.339500e-03
treatment_type_Surgery	0.002839	6.018004e-05
treatment_type_Unknown	-0.003513	-9.354595e-19

	cholesterol_level	hypertension	asthma
age	2.242015e-03	0.000134	-0.002869
gender	2.760413e-03	0.002087	-0.001366
cancer_stage	-2.119179e-03	0.008178	-0.005568
family_history	-3.579847e-04	-0.004445	-0.005626
bmi	7.480522e-01	0.002758	-0.000306
cholesterol_level	1.000000e+00	0.000446	-0.001701
hypertension	4.455175e-04	1.000000	0.103190
asthma	-1.700743e-03	0.103190	1.000000
cirrhosis	-1.288772e-03	0.097863	0.055177
other_cancer	-3.940041e-03	0.074388	0.036185
survived	-2.410907e-03	-0.001534	0.003804
treatment_time	-5.801593e-03	-0.013976	-0.001586
smoking_status_Former Smoker	-6.157001e-03	-0.000170	0.004439

```
plt.figure(figsize=(20,15))
sns.heatmap(corr, annot=True)
plt.show()
```



```
category = ['gender', 'country', 'cancer_stage',
            'family_history']

for i in category:
    print(f"{i.capitalize()}\nUnique values: {df_model[i].unique()}\nCount of uni
```

```
Gender
Unique values: [0 1]
Count of unique values: 2

Country
Unique values: ['Sweden' 'Netherlands' 'Hungary' 'Belgium' 'Luxembourg' 'It
'Denmark' 'Malta' 'Germany' 'Poland' 'Ireland' 'Romania' 'Spain' 'Greece'
'Estonia' 'Cyprus' 'France' 'Slovenia' 'Latvia' 'Portugal' 'Austria'
'Czech Republic' 'Finland' 'Lithuania' 'Slovakia' 'Bulgaria']
Count of unique values: 27

Cancer_stage
Unique values: [1 3 4 2 0]
Count of unique values: 5

Family_history
Unique values: [1 0]
Count of unique values: 2
```

```
one_hot_columns = [
    'gender',
    'family_history',
]

df_encoded = pd.get_dummies(df_model[one_hot_columns], drop_first=False)

df_encoded = df_encoded.replace({True: 1, False: 0}).astype(int)
```

```

from sklearn.model_selection import train_test_split

# Define features (X) and target (y)
X = df_model.drop('survived', axis=1)
y = df_model['survived']

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random

print("Training set shape:", X_train.shape)
print("Testing set shape:", X_test.shape)

```

```

↗ Training set shape: (64586, 20)
  Testing set shape: (16147, 20)

```

```

from imblearn.over_sampling import SMOTE

# Apply SMOTE only to the training data
smote = SMOTE(random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train.select_dtypes(include

print("Shape of training data after SMOTE:", X_train_smote.shape)
print("Distribution of 'survived' in training data after SMOTE:")
print(y_train_smote.value_counts())

```

```

↗ Shape of training data after SMOTE: (100624, 11)
  Distribution of 'survived' in training data after SMOTE:
  survived
0      50312
1      50312
Name: count, dtype: int64

```

```

# Initialize models
log_reg = LogisticRegression(max_iter=1000, random_state=42)
rf_clf = RandomForestClassifier(random_state=42)
xgb_clf = XGBClassifier(random_state=42)

models = {
    "Logistic Regression": log_reg,
    "Random Forest": rf_clf,
    "XGBoost": xgb_clf
}

accuracy_scores = {}

# Train and evaluate models
for name, model in models.items():

```

```

print(f"Training {name}...")
model.fit(X_train_smote, y_train_smote)
y_pred = model.predict(X_test.select_dtypes(include=np.number))
report = classification_report(y_test, y_pred, output_dict=True)
accuracy = report['accuracy']
accuracy_scores[name] = accuracy
print(f"Evaluation for {name}:")
print(classification_report(y_test, y_pred))
print("-" * 30)

```



Training Logistic Regression...

Evaluation for Logistic Regression:

	precision	recall	f1-score	support
0	0.77	0.38	0.51	12578
1	0.21	0.60	0.32	3569
accuracy			0.43	16147
macro avg	0.49	0.49	0.41	16147
weighted avg	0.65	0.43	0.46	16147

-----

Training Random Forest...

Evaluation for Random Forest:

	precision	recall	f1-score	support
0	0.78	0.85	0.82	12578
1	0.23	0.16	0.19	3569
accuracy			0.70	16147
macro avg	0.51	0.51	0.50	16147
weighted avg	0.66	0.70	0.68	16147

-----

Training XGBoost...

Evaluation for XGBoost:

	precision	recall	f1-score	support
0	0.78	0.95	0.86	12578
1	0.24	0.06	0.09	3569
accuracy			0.75	16147
macro avg	0.51	0.50	0.47	16147
weighted avg	0.66	0.75	0.69	16147

-----

```
# Plot the accuracy of the models
plt.figure(figsize=(10, 6))
sns.barplot(x=list(accuracy_scores.keys()), y=list(accuracy_scores.values()), palette="magma")
plt.title("Accuracy of Different Models")
plt.ylabel("Accuracy")
plt.ylim(0, 1)
plt.show()
```

